

Nettverksprogrammering

Atomic typer, parallellisering og prosesser

Ole Christian Eidheim

Institutt for datateknologi og informatikk,
NTNU

15. januar 2024

Oversikt

- **Atomic-typer**
- CPU/GPU parallellisering
- Prosesser

Atomic-typer

- trenger ikke å bli beskyttet av mutex

Ulempe: bare enkle datatyper som int og float kan gjøres atomic.

```
#include <atomic>
#include <iostream>
#include <thread>

using namespace std;

int main() {
    atomic<int> sum(0);

    thread t1([&sum]() {
        for (int c = 0; c < 1000; c++)
            sum++;
    });
    thread t2([&sum]() {
        for (int c = 0; c < 1000; c++)
            sum++;
    });
    t1.join();
    t2.join();

    cout << sum << endl; // Output: 2000
}
```

Atomic-typer

- Handtering av tilstand i flere tråder

```
#include <iostream>
#include <thread>

using namespace std;

enum class State { sitting, standing_up, standing };

int main() {
    atomic<State> state(State::sitting);

    thread([&state] { /* Draw animation frames based on state */ });

    while (true) { // Handle input
        if (/* keypress */) {
            // Stand up if sitting:
            auto expected = State::sitting;
            if (state.compare_exchange_strong(expected, State::standing_up)) {
                // Standing up, play squeaky chair sound
            }
        }
    }
}
```

Legg merke til den spesielle funksjonen `compare_exchange_strong()`, og at vi slipper å bruke mutex selv om vi både leser og skriver til state.

Mer om atomic-typer

- referansetelling

- En form for *garbage collection* der et objekt blir frigjort når det ikke lenger blir brukt
- Kan være nyttig i trådprogrammering der en ikke vet i hvilken tråd et objekt brukes for siste gang

Eksempler:

- C++: `std::shared_ptr`
- Rust: `std::sync::Arc` (Atomically Reference Counted)

```
#include <iostream>
#include <thread>

using namespace std;

int main() {
    thread t;
    {
        std::shared_ptr<int> ref_counted(new int(42));
        t = thread([ref_counted] { // ref_counted is copied to thread
            this_thread::sleep_for(1s); // Wait 1 second
            cout << "value from thread: " << *ref_counted << endl;
            cout << "count from thread: " << ref_counted.use_count() << endl;
            // The last ref_counted object is destroyed at end of thread,
            // and its int value is then freed from memory
        });
        cout << "value: " << *ref_counted << endl;
        cout << "count: " << ref_counted.use_count() << endl;
        // One ref_counted object is destroyed at end of scope,
        // and its use_count is reduced by 1
    }
    t.join();
}

// Output:
// value: 42
// count: 2
// value from thread: 42
// count from thread: 1
```

Mer om atomic-typer

- referansetelling implementasjon, men ikke trådsikker

```
#include <iostream>
using namespace std;

class RefCountedInt {
public:
    class Object {
    public:
        int *ptr;
        int count;
        Object(int *ptr_) : ptr(ptr_), count(1) {}
        ~Object() { delete ptr; }
    };
    Object *object;

    RefCountedInt(int *ptr) : object(new Object(ptr)) {
        cout << "constructor" << endl;
    }
    ~RefCountedInt() {
        cout << "destructor" << endl;
        object->count--;
        if (object->count == 0) {
            cout << "deleting object" << endl;
            delete object;
        }
    }
    RefCountedInt(const RefCountedInt &other) {
        cout << "copy constructor" << endl;
        object = other.object;
        object->count++;
    }
};
```

```
void f(RefCountedInt ref_counted) {
    cout << "count: " << ref_counted.object->count << endl;
}

int main() {
    RefCountedInt ref_counted(new int(42));
    auto a = ref_counted;
    f(a);
}

// Output:
// constructor
// copy constructor
// copy constructor
// count: 3
// destructor
// destructor
// destructor
// deleting object
```

Ikke så farlig om du ikke forstår all koden i venstre kolonne, men se forskjellene på neste slide.

Mer om atomic-typer

- trådsikker referansetelling implementasjon med atomic

```
#include <iostream>
using namespace std;

class RefCountedInt {
public:
    class Object {
    public:
        int *ptr;
        atomic<int> count;
        Object(int *ptr_) : ptr(ptr_), count(1) {}
        ~Object() { delete ptr; }
    };
    Object *object;

    RefCountedInt(int *ptr) : object(new Object(ptr)) {
        cout << "constructor" << endl;
    }
    ~RefCountedInt() {
        cout << "destructor" << endl;
        auto previous_count = object->count.fetch_sub(1);
        if (previous_count == 1) {
            cout << "deleting object" << endl;
            delete object;
        }
    }
    RefCountedInt(const RefCountedInt &other) {
        cout << "copy constructor" << endl;
        object = other.object;
        object->count++;
    }
};
```

```
void f(RefCountedInt ref_counted) {
    cout << "count: " << ref_counted.object->count << endl;
}

int main() {
    RefCountedInt ref_counted(new int(42));
    auto a = ref_counted;
    f(a);
}

// Output:
// constructor
// copy constructor
// copy constructor
// count: 3
// destructor
// destructor
// destructor
// deleting object
```


Legg merke til den spesielle funksjonen `fetch_sub()`.

Oversikt


- Atomic-typer
- **CPU/GPU** parallellisering
- Prosesser

Parallellisering CPU vs GPU

CPU (Multiple Cores)



GPU (Hundreds of Cores)



CPU parallellisering

- skal parallellisere dette

```
#include <iostream>
#include <vector>

using namespace std;

int main() {
    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    for (int i = 0; i < 10; i++) {
        c[i] = a[i] + b[i];
    }

    // c: 0 2 4 3 5 7 6 8 10 9
}
```

CPU parallellisering

- manuell tungvint løsning, kun CPU

```
#include <iostream>
#include <thread>
#include <vector>

using namespace std;

int main() {
    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    vector<thread> threads;
    for (int thread_number = 0; thread_number < 5; thread_number++) {
        threads.emplace_back([thread_number, &a, &b, &c] {
            for (int i = thread_number * 2; i <= thread_number * 2 + 1; i++)
                c[i] = a[i] + b[i];
        });
    }

    for (auto &t : threads)
        t.join();

    // c: 0 2 4 3 5 7 6 8 10 9
}
```

Suboptimal CPU parallelising

- OpenMP (Open Multi-Processing)

```
#include <iostream>
#include <vector>

using namespace std;

int main() {
    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    #pragma omp parallel for
    for (int i = 0; i < 10; i++) {
        c[i] = a[i] + b[i];
    }

    // c: 0 2 4 3 5 7 6 8 10 9
}

// Compile with g++ and add the flag -fopenmp
```

CPU parallellisering

- std::algorithm før c++17 ingen parallellisering

```
#include <algorithm>
#include <iostream>
#include <vector>

using namespace std;

int main() {
    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    transform(a.begin(), a.end(), b.begin(), c.begin(),
        [](int a_element, int b_element) {
            return a_element + b_element;
        });

    // c: 0 2 4 3 5 7 6 8 10 9
}
```

CPU(/fremtidig GPU?) parallellisering

- std::algorithm c++17

```
#include <algorithm>
#include <execution>
#include <iostream>
#include <vector>

using namespace std;

int main() {
    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    transform(execution::par, a.begin(), a.end(), b.begin(), c.begin(),
              [](int a_element, int b_element) {
                  return a_element + b_element;
              });

    // c: 0 2 4 3 5 7 6 8 10 9
}

// Compile using a newer g++ version with the flags: -ltbb -std=c++17
```

GPU parallellisering

- Komplisert OpenCL (Open Computing Language) kode

```
int a[10] = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
int b[10] = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
int c[10];

string kernel_code =
    "void kernel simple_add(global const int* a, global const int* b, global int* c) {"
    "    c[get_global_id(0)] = a[get_global_id(0)] + b[get_global_id(0)];"
    "}";

cl::Program program(/*her velges kernel_code og OpenCL parametere*/);
cl::Kernel kernel_add = cl::Kernel(program, "simple_add");

cl::CommandQueue queue(/*her settes OpenCL parametere*/);
cl::Buffer device_a(/*OpenCL parameter*/, CL_MEM_READ_WRITE, sizeof(int) * 10);
cl::Buffer device_b(/*OpenCL parameter*/, CL_MEM_READ_WRITE, sizeof(int) * 10);
cl::Buffer device_c(/*OpenCL parameter*/, CL_MEM_READ_WRITE, sizeof(int) * 10);
queue.enqueueWriteBuffer(device_a, CL_TRUE, 0, sizeof(int) * 10, a);
queue.enqueueWriteBuffer(device_b, CL_TRUE, 0, sizeof(int) * 10, b);
kernel_add.setArg(0, device_a);
kernel_add.setArg(1, device_b);
kernel_add.setArg(2, device_c);

//Programmet kjøres på den valgte enheten (feks GPU):
queue.enqueueNDRangeKernel(kernel_add, cl::NullRange, cl::NDRange(10), cl::NullRange);
queue.finish();
queue.enqueueReadBuffer(device_c, CL_TRUE, 0, sizeof(int) * 10, c);

// c: 0 2 4 3 5 7 6 8 10 9
```

GPU parallellisering

- Boost.Compute (se eksempelet [her](#))

```
#include <boost/compute/algorithm/transform.hpp>
#include <boost/compute/container/vector.hpp>
#include <iostream>

using namespace std;
namespace compute = boost::compute;

int main() {
    auto device = compute::system::default_device();
    compute::context context(device);
    compute::command_queue queue(context, device);

    vector<int> a = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
    vector<int> b = {0, 1, 2, 0, 1, 2, 0, 1, 2, 0};
    vector<int> c(10);

    compute::vector<int> device_a(a.size(), context);
    compute::vector<int> device_b(b.size(), context);
    compute::copy(a.begin(), a.end(), device_a.begin(), queue);
    compute::copy(b.begin(), b.end(), device_b.begin(), queue);

    compute::vector<int> device_c(c.size(), context);
    compute::transform(device_a.begin(), device_a.end(),
                       device_b.begin(), device_c.begin(), compute::plus<int>(), queue);
    compute::copy(device_c.begin(), device_c.end(), c.begin(), queue);

    // c: 0 2 4 3 5 7 6 8 10 9
}
```


Andre CPU/GPU parallelliseringsbiblioteker

■ ArrayFire

- Startet i 2014
- C++ bibliotek
- Støtter CUDA, OpenCL og CPU
- Kan enkelt installeres med for eksempel pacman (Arch Linux / Manjaro) eller brew (MacOS)

■ Kompute

- Startet i 2020
- C++ bibliotek
- Støtter CUDA, OpenCL og Vulkan.

- 2021: Mangler dessverre enda slike biblioteker i Rust

Oversikt

- Atomic-typer
- CPU/GPU parallellisering
- **Prosesser**

Prosesser

- Tråder
 - Kjører i delt minneområde
 - Programmeringsfeil kan føre til at en tråd får tilgang til minneområdet til en annen tråd
 - Krasj i en tråd krasjer hele programmet
- Prosesser
 - Kjører i separate minneområder
 - Sikrere mot programmeringsfeil, men kommunikasjon mellom prosesser er mer ressurskrevende
 - En krasj vil ikke påvirke andre prosesser

Prossesser, lese data fra prosess

- eksempler med `tiny-process-library`

```
#include "process.hpp"
#include <iostream>

using namespace std;
using namespace TinyProcessLib;

int main() {
    Process process("echo Hello World", {},
        [](const char *bytes, size_t n) {
            cout << string(bytes, n); // Output: Hello World
        });

    cout << process.get_exit_status() << endl; // Output: 0
}
```

Prossesser, lese data fra prosess

- eksempler med [tiny-process-library](#)

```
#include "process.hpp"
#include <iostream>

using namespace std;
using namespace TinyProcessLib;

int main() {
    Process process("cat nonexistent_file", {},
        [](const char *bytes, size_t n) {
            cout << string(bytes, n); // No output
        }, [](const char *bytes, size_t n) {
            // Output: nonexistent_file: No such file or directory
            cout << string(bytes, n);
        });

    cout << process.get_exit_status() << endl; // Output: 1
}
```

Prossesser, lese data fra *inline* prosess

- eksempler med [tiny-process-library](#)

```
#include "process.hpp"
#include <iostream>

using namespace std;
using namespace TinyProcessLib;

int main() {
    Process process([] { // Does not work on Windows
                        // where an executable file is needed
        cout << "Hello" << endl;
        cerr << "World" << endl;
        exit(10);
    }, [](const char *bytes, size_t n) {
        cout << string(bytes, n); // Output: Hello
    }, [](const char *bytes, size_t n) {
        cout << string(bytes, n); // Output: World
    });

    cout << process.get_exit_status() << endl; // Output: 10
}
```

Prossesser, *inline* prosess

- eksempler med *tiny-process-library*

Hva skjer her?

```
#include "process.hpp"
#include <iostream>

using namespace std;
using namespace TinyProcessLib;

int main() {
    int a = 42;

    Process process([&a] { // Does not work on Windows
                                // where an executable file is needed

        a++;
        cout << a << endl;
        exit(0);
    }, [](const char *data, size_t n) {
        cout << string(data, n); // Output: 43
    });

    cout << process.get_exit_status() << endl; // Output: 0

    cout << a << endl; // Output 42
}
```

Prosesser, skrive til og lese data fra prosess

- eksempler med [tiny-process-library](#)

```
#include "process.hpp"
#include <iostream>

using namespace std;
using namespace TinyProcessLib;

int main() {
    Process process("cat", {},
        [](const char *bytes, size_t n) {
            cout << string(bytes, n); // Output: Hello World
        }, nullptr /* no stderr */, true /* open stdin */);

    process.write("Hello World\n");
    process.close_stdin();

    cout << process.get_exit_status() << endl; // Output: 0
}
```